

- **Vendor: Microsoft**
- **Exam Code: AZ-204**
- **Exam Name: Developing Solutions for Microsoft Azure**
- **Part of New Questions from [PassLeader](#) (Updated in [Nov/2020](#))**

[Visit PassLeader and Download Full Version AZ-204 Exam Dumps](#)

NEW QUESTION 280

You develop a website. You plan to host the website in Azure. You expect the website to experience high traffic volumes after it is published. You must ensure that the website remains available and responsive while minimizing cost. You need to deploy the website. What should you do?

- A. Deploy the website to a virtual machine.
Configure the virtual machine to automatically scale when the CPU load is high.
- B. Deploy the website to an App Service that uses the Shared service tier.
Configure the App Service plan to automatically scale when the CPU load is high.
- C. Deploy the website to a virtual machine.
Configure a Scale Set to increase the virtual machine instance count when the CPU load is high.
- D. Deploy the website to an App Service that uses the Standard service tier.
Configure the App Service plan to automatically scale when the CPU load is high.

Answer: D

Explanation:

Windows Azure Web Sites (WAWS) offers 3 modes: Standard, Free, and Shared. Standard mode carries an enterprise-grade SLA (Service Level Agreement) of 99.9% monthly, even for sites with just one instance. Standard mode runs on dedicated instances, making it different from the other ways to buy Windows Azure Web Sites.

Incorrect:

Not B: Shared and Free modes do not offer the scaling flexibility of Standard, and they have some important limits. Shared mode, just as the name states, also uses shared Compute resources, and also has a CPU limit. So, while neither Free nor Shared is likely to be the best choice for your production environment due to these limits.

NEW QUESTION 281

You develop an HTTP triggered Azure Function app to process Azure Storage blob data. The app is triggered using an output binding on the blob. The app continues to time out after four minutes. The app must process the blob data. You need to ensure the app does not time out and processes the blob data.

Solution: Use the Durable Function async pattern to process the blob data.

Does the solution meet the goal?

- A. Yes
- B. No

Answer: B

Explanation:

Instead pass the HTTP trigger payload into an Azure Service Bus queue to be processed by a queue trigger function and return an immediate HTTP success response.

<https://docs.microsoft.com/en-us/azure/azure-functions/functions-best-practices>

NEW QUESTION 282

You develop an HTTP triggered Azure Function app to process Azure Storage blob data. The app is triggered using an output binding on the blob. The app continues to time out after four minutes. The app must process the blob data. You need to ensure the app does not time out and processes the blob data.

Solution: Pass the HTTP trigger payload into an Azure Service Bus queue to be processed by a queue trigger function and return an immediate HTTP success response.

Does the solution meet the goal?

- A. Yes
- B. No

Answer: A

Explanation:

Large, long-running functions can cause unexpected timeout issues. General best practices include: whenever possible, refactor large functions into smaller function sets that work together and return responses fast. For example, a webhook or HTTP trigger function might require an acknowledgment response within a certain time limit; it's common for webhooks to require an immediate response. You can pass the HTTP trigger payload into a queue to be processed by a queue trigger function. This approach lets you defer the actual work and return an immediate response.

<https://docs.microsoft.com/en-us/azure/azure-functions/functions-best-practices>

NEW QUESTION 283

You are developing an Azure Cosmos DB solution by using the Azure Cosmos DB SQL API. The data includes millions of documents. Each document may contain hundreds of properties. The properties of the documents do not contain distinct values for partitioning. Azure Cosmos DB must scale individual containers in the database to meet the performance needs of the application by spreading the workload evenly across all partitions over time. You need to select a partition key. Which two partition keys can you use? (Each correct answer presents a complete solution. Choose two.)

- A. a single property value that does not appear frequently in the documents
- B. a value containing the collection name
- C. a single property value that appears frequently in the documents
- D. a concatenation of multiple property values with a random suffix appended
- E. a hash suffix appended to a property value

Answer: DE

Explanation:

You can form a partition key by concatenating multiple property values into a single artificial partitionKey property. These keys are referred to as synthetic keys. Another possible strategy to distribute the workload more evenly is to append a random number at the end of the partition key value. When you distribute items in this way, you can perform parallel write operations across partitions.

<https://docs.microsoft.com/en-us/azure/cosmos-db/synthetic-partition-keys>

NEW QUESTION 284

You are building a website that uses Azure Blob storage for data storage. You configure Azure Blob storage lifecycle to move all blobs to the archive tier after 30 days. Customers have requested

a service-level agreement (SLA) for viewing data older than 30 days. You need to document the minimum SLA for data recovery. Which SLA should you use?

- A. at least two days
- B. between one and 15 hours
- C. at least one day
- D. between zero and 60 minutes

Answer: B

Explanation:

The archive access tier has the lowest storage cost. But it has higher data retrieval costs compared to the hot and cool tiers. Data in the archive tier can take several hours to retrieve depending on the priority of the rehydration. For small objects, a high priority rehydrate may retrieve the object from archive in under 1 hour.

<https://docs.microsoft.com/en-us/azure/storage/blobs/storage-blob-storage-tiers?tabs=azure-portal>

NEW QUESTION 285

You develop an app that allows users to upload photos and videos to Azure storage. The app uses a storage REST API call to upload the media to a blob storage account named Account1. You have blob storage containers named Container1 and Container2. Uploading of videos occurs on an irregular basis. You need to copy specific blobs from Container1 to Container2 when a new video is uploaded. What should you do?

- A. Copy blobs to Container2 by using the Put Blob operation of the Blob Service REST API.
- B. Create an Event Grid topic that uses the Start-AzureStorageBlobCopy cmdlet.
- C. Use AzCopy with the Snapshot switch to copy blobs to Container2.
- D. Download the blob to a virtual machine and then upload the blob to Container2.

Answer: B

Explanation:

The Start-AzureStorageBlobCopy cmdlet starts to copy a blob.

<https://docs.microsoft.com/en-us/powershell/module/azure.storage/start-azurestorageblobcopy?view=azurermps-6.13.0>

NEW QUESTION 286

You are developing an ASP.NET Core website that uses Azure FrontDoor. The website is used to build custom weather data sets for researchers. Data sets are downloaded by users as Comma Separated Value (CSV) files. The data is refreshed every 10 hours. Specific files must be purged from the FrontDoor cache based upon Response Header values. You need to purge individual assets from the Front Door cache. Which type of cache purge should you use?

- A. single path
- B. wildcard
- C. root domain

Answer: A

Explanation:

<https://docs.microsoft.com/en-us/azure/frontdoor/front-door-caching>

NEW QUESTION 287

You are developing a Java application that uses Cassandra to store key and value data. You plan to use a new Azure Cosmos DB resource and the Cassandra API in the application. You create an Azure Active Directory (Azure AD) group named Cosmos DB Creators to enable provisioning of Azure Cosmos accounts, databases, and containers. The Azure AD group must not be able to

access the keys that are required to access the data. You need to restrict access to the Azure AD group. Which role-based access control should you use?

- A. DocumentDB Accounts Contributor
- B. Cosmos Backup Operator
- C. Cosmos DB Operator
- D. Cosmos DB Account Reader

Answer: C

Explanation:

Azure Cosmos DB now provides a new RBAC role, Cosmos DB Operator. This new role lets you provision Azure Cosmos accounts, databases, and containers, but can't access the keys that are required to access the data. This role is intended for use in scenarios where the ability to grant access to Azure Active Directory service principals to manage deployment operations for Cosmos DB is needed, including the account, database, and containers.

<https://azure.microsoft.com/en-us/updates/azure-cosmos-db-operator-role-for-role-based-access-control-rbac-is-now-available/>

NEW QUESTION 288

You develop and deploy an ASP.NET web app to Azure App Service. You use Application Insights telemetry to monitor the app. You must test the app to ensure that the app is available and responsive from various points around the world and at regular intervals. If the app is not responding, you must send an alert to support staff. You need to configure a test for the web app. Which two test types can you use? (Each correct answer presents a complete solution. Choose two.)

- A. integration
- B. multi-step web
- C. URL ping
- D. unit
- E. load

Answer: BC

Explanation:

There are three types of availability tests:

- URL ping test: a simple test that you can create in the Azure portal.
 - Multi-step web test: A recording of a sequence of web requests, which can be played back to test more complex scenarios. Multi-step web tests are created in Visual Studio Enterprise and uploaded to the portal for execution.
 - Custom Track Availability Tests: If you decide to create a custom application to run availability tests, the TrackAvailability() method can be used to send the results to Application Insights.
- <https://docs.microsoft.com/en-us/azure/azure-monitor/app/monitor-web-app-availability>

NEW QUESTION 289

You are developing an e-commerce solution that uses a microservice architecture. You need to design a communication backplane for communicating transactional messages between various parts of the solution. Messages must be communicated in first-in-first-out (FIFO) order. What should you use?

- A. Azure Storage Queue
- B. Azure Event Hub
- C. Azure Service Bus
- D. Azure Event Grid

Answer: A

Explanation:

As a solution architect/developer, you should consider using Service Bus queues when: your solution requires the queue to provide a guaranteed first-in-first-out (FIFO) ordered delivery.



<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-azure-and-service-bus-queues-compared-contrasted>

NEW QUESTION 290

Hotspot

A company is developing a gaming platform. Users can join teams to play online and see leaderboards that include player statistics. The solution includes an entity named Team. You plan to implement an Azure Redis Cache instance to improve the efficiency of data operations for entities that rarely change. You need to invalidate the cache when team data is changed. How should you complete the code? (To answer, select the appropriate options in the answer area.)

Answer Area

```
void ClearCachedTeams()
{
    
    
    ViewBag.msg += "Team data removed from cache.";
}
```

Answer:

Answer Area

```
void ClearCachedTeams()
{
    

IDatabase cache = Connection.GetDatabase();  
ICache cache = Connection.GetDatabase();



cache.KeyDelete("Team");  
cache.StringSet("Team", "");  
cache.ValueDelete("Team");  
cache.StringGet("Team", "");


    ViewBag.msg += "Team data removed from cache.";
}
```

Explanation:

Box 1: Connection refers to a previously configured ConnectionMultiplexer.

Box 2: To specify the expiration of an item in the cache, use the TimeSpan parameter of StringSet.cache.StringSet("key1", "value1", TimeSpan.FromMinutes(90));.

<https://azure.microsoft.com/sv-se/blog/lap-around-azure-redis-cache-preview/>

<https://docs.microsoft.com/en-us/cli/azure/webapp/config/container>

NEW QUESTION 291

Drag and Drop

Your company has several websites that use a company logo image. You use Azure Content Delivery Network (CDN) to store the static image. You need to determine the correct process of how the CDN and the Point of Presence (POP) server will distribute the image and list the items in the correct order. In which order do the actions occur? (To answer, move all actions from the list of actions to the answer area and arrange them in the correct order.)

Answer Area

The origin server returns the logo image to an edge server in the POP. An edge server in the POP caches the logo image and returns the image to the client.



Actions

Subsequent requests for the file may be directed to the same POP using the CDN logo image URL. The POP edge server returns the file from cache if the TTL has not expired.



Step 4: Additional users can then request the same file by using the same URL that the original user used, and can also be directed to the same POP. If the TTL for the file hasn't expired, the POP edge server returns the file directly from the cache. This process results in a faster, more responsive user experience.

<https://docs.microsoft.com/en-us/azure/cdn/cdn-overview>

NEW QUESTION 292

.....

[Visit PassLeader and Download Full Version AZ-204 Exam Dumps](#)